

Course Description

- Covers techniques for managing data throughout an end-to-end ML process
- Learn statistical analysis and visualization techniques, including methods for detecting and remedying overfitting
- Gain familiarity with tools in standard ML and data processing libraries

Course Description

- **Lecture 1:** ML Life-cycle, Data Analysis and Basic ML model: K Nearest Neighbors
- **Lecture 2:** Feature Engineering, Decision Trees, Hyper-parameters and AWS Sagemaker
- **Lecture 3:** Optimization, Regression Models, Boosting and Neural Networks


Final Project-Completion

- Apply and experiment with a real-world Amazon dataset.
- Completion of this course is based on your Leaderboard submission.
- Submission page:
<https://leaderboard.corp.amazon.com/tasks/478>
- Submission is open until Saturday 5:00 PM (PST)

Top 3 submissions - this week



Completions



Final Project-Completion

You will get a score after each submission. You can improve your score by making multiple submissions (no upper limit on number of submissions).

- **Requirements:** At least **ONE** submission is **REQUIRED**
- **Non-completion:** Submit 0 models to the Class Leaderboard



After completion, student and manager receive confirmation email.



Topics for today

- Optimization
- Regression
- Boosting
- Neural Networks
- MxNet
- Final Project

Topics for today

- Optimization
- Regression
- Boosting
- Neural Networks
- MxNet
- Final Project

Optimization

- Optimization is an important part of our life.
- For example we:
 - Optimize the order of our work tasks based on priorities.
 - Optimize our route home to minimize our travel time.
- In machine learning, similarly, we use optimization to usually minimize some error rate of our ML model.
- We will learn the Gradient Descent Optimization.

Gradient Descent


- In machine learning, we usually optimize functions.

$y = f(x)$ where y =output, x =input, f =function

- **Minimizing f** means **finding** the input x that results in **the lowest** value y .
(Maximize: Find x that gives the largest y)
- **Gradient:**
 - **Definition:** It is the direction and rate of the fastest increase in a function.
 - It can be calculated with **partial derivatives** of the function f with respect to each input variable in x .
 - Because it has a direction and rate, we also call it a “vector”.

Gradient Example


$$y = f(x) = 0.6 x^2$$



$$y = f(x) = 0.6 x^2 \text{ and gradient is } \nabla f = \langle 1.2x \rangle$$




Gradient Example



$y = f(x) = 0.6 x^2$ and gradient is $\nabla f = \langle 1.2x \rangle$

$$\nabla f(2) = \langle 2.4 \rangle$$

Gradient Example




$y = f(x) = 0.6 x^2$ and gradient is $\nabla f = \langle 1.2x \rangle$

$$\nabla f(2) = \langle 2.4 \rangle$$

$$\nabla f(4) = \langle 4.8 \rangle$$

Gradient Example




$y = f(x) = 0.6 x^2$ and gradient is $\nabla f = \langle 1.2x \rangle$

$$\nabla f(2) = \langle 2.4 \rangle$$

$$\nabla f(4) = \langle 4.8 \rangle$$

$$\nabla f(-3) = \langle -3.6 \rangle$$

Gradient Example



$y = f(x) = 0.6 x^2$ and gradient is $\nabla f = \langle 1.2x \rangle$


$$\nabla f(2) = \langle 2.4 \rangle$$

$$\nabla f(4) = \langle 4.8 \rangle$$

$$\nabla f(-3) = \langle -3.6 \rangle$$

$$\nabla f(0) = \langle 0 \rangle$$

Gradient Example



$y = f(x) = 0.6 x^2$ and gradient is $\nabla f = \langle 1.2x \rangle$

$$\nabla f(2) = \langle 2.4 \rangle \Rightarrow |\nabla f(2)| = 2.4$$

$$\nabla f(4) = \langle 4.8 \rangle \Rightarrow |\nabla f(4)| = 4.8$$

$$\nabla f(-3) = \langle -3.6 \rangle \Rightarrow |\nabla f(-3)| = 3.6$$

$$\nabla f(0) = \langle 0 \rangle \Rightarrow |\nabla f(0)| = 0$$

- Sign of the gradient shows direction: + right and – left (the direction the func. increases)
- As we go towards to the bottom part of the function, gradient length gets smaller and becomes zero

Gradient - Math

Writing the formula:

$$\nabla f = \text{grad } f = \left\langle \frac{\partial f}{\partial x}(x) \right\rangle$$

$\langle \rangle$ shows that gradient is a "vector" here and $\frac{\partial y}{\partial x}$ is used for partial derivative

Example:

$$y = f(x) = 0.6 x^2$$

$$\text{Gradient is } \nabla f = \langle 1.2x \rangle$$

Gradient – Math (Higher Dim.)

Writing the formula:

$$\nabla f = \text{grad } f = \left\langle \frac{\partial f}{\partial x_1}(x_1, x_2, \dots, x_n), \frac{\partial f}{\partial x_2}(x_1, x_2, \dots, x_n), \dots, \frac{\partial f}{\partial x_n}(x_1, x_2, \dots, x_n) \right\rangle$$

$\langle \rangle$ shows that gradient is a "vector" here and $\frac{\partial y}{\partial x}$ is used for partial derivative


Example:

$$y = f(x_1, x_2) = 3x_1^2 + 5x_2^2$$


$$\text{Gradient is } \nabla f = \langle 6x_1, 10x_2 \rangle$$

Gradient – Math (Higher Dim.)

$y = f(x_1, x_2) = 3x_1^2 + 5x_2^2$ and gradient is $\nabla f = \langle 6x_1, 10x_2 \rangle$



$$\nabla f(3, 2) = \langle 18, 20 \rangle$$



Gradient Descent Method

- Gradient descent method uses gradients to find the **minimum** of a function **iteratively**.
- We take **small steps** towards the minimum (opposite direction of gradient).

Starting at an initial point x


- Update x_{new} : $x_{\text{current}} - \text{step_size} * \text{gradient}$
- Let's find the minimum of $y=0.6x^2$, we will start from an initial point assume $x=4.25$ and $\text{step_size}=0.2$

Gradient Descent Method

- Let's find the minimum of $y=0.6x^2$, we will start from an initial point assume $x=4.25$ and $\text{step_size}=0.2$

Gradient is $\nabla f = \langle 1.2x \rangle$


$x_{\text{new}}: x_{\text{curr.}} - \text{step_size} * \text{gradient}$



Topics for today

- Optimization
- Regression
- Boosting
- Neural Networks
- MxNet
- Final Project


Linear Regression



- We use it for numerical value prediction.
- For example: Predicting house prices by looking at square footage, number of rooms etc.
- When there is a single feature (explanatory variable x) and a real-valued response (target variable y):

$$y = w_0 + w_1x$$

Linear Regression



- Given data (x, y) , a line: $\mathbf{y} = \mathbf{w}_0 + \mathbf{w}_1x$ is defined by w_0 (i.e. intercept) and w_1 (slope)
- The vertical offset for each data point from the line is the error between the true label y and the prediction based on x .
- The best line minimizes the sum of squared errors (SSE): $\sum (y_i - \hat{y}_i)^2$

Linear Regression

- Multiple linear regression includes m features with $m \geq 2$:

$$y = w_0 + w_1x_1 + \dots + w_mx_m$$

- Sensitive to correlation between features resulting in high variance of coefficients.
- Features are allowed to have interactions and higher order terms.

Gradient Descent - Fitting a model

- Let's have a linear regression model:

$$y = w_0 + w_1x_1 + w_2x_2 + \dots + w_qx_q$$

where w_0, w_1, \dots, w_m : coefficients and x_0, x_1, \dots, x_m : input variables

- Minimize the cost function Mean Squared Error:

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 \text{ where } \mathbf{y}_i: \text{Data value, } \hat{\mathbf{y}}_i: \text{Predicted value and } n: \# \text{ of records}$$

- Iteratively update coefficients with Gradient Descent:

$$w_{\text{new}} = w_{\text{current}} - \text{step_size} * \text{gradient}$$

Logistic Regression

Linear regression was useful when predicting continuous values.

$$y = \mathbf{w_0} + \mathbf{w_1x_1} + \mathbf{w_2x_2} + \cdots + \mathbf{w_qx_q}$$

Can we use a similar approach to solve classification problems?

The most simple classification problem is binary classification $y=0$ or 1 :


Examples:

- Email: Spam or Not spam
- Text: Positive or negative meaning
- Image: Hotdog or Not hotdog

How can we modify this linear regression equation to work for this?

$$y = \mathbf{w_0} + \mathbf{w_1x_1} + \mathbf{w_2x_2} + \cdots + \mathbf{w_qx_q}$$

Applying Sigmoid Function



- $\text{sigmoid}(x) = \frac{1}{1+e^{-x}}$
- We can classify with this:
 - if $\text{sigmoid}(x) < 0.5$, round down:(0)
 - if $\text{sigmoid}(x) \geq 0.5$, round up:(1)
- Our regression function becomes:

$$h_w(x) = p = \frac{1}{1 + e^{-(w_0 + w_1x_1 + w_2x_2 + \dots + w_qx_q)}}$$

We can say p is the probability that $y=1$ for input x .

Classify as 1 if $p \geq 0.5$, as 0 if $p < 0.5$

Log-loss (Binary Cross-entropy)

Log-loss: A numeric value that measures the performance of a binary classifier when the output of the model is probability between 0 and 1.

- We prefer small values for the loss (we try to minimize it). A loss of 0 means the perfect classifier.
- Log loss gets larger when the predicted probability diverges from the actual label (0 or 1).
- In mathematical terms:

$$\text{Logloss} = -(y * \log(p) + (1 - y) * \log(1 - p))$$

where


y: True class, p: Prob. of the model, log: Logarithm function

Log-loss (Binary Cross-entropy)

Example: Let's calculate the log-loss for these situations below

$$\text{Logloss} = -(y * \log(p) + (1 - y) * \log(1 - p))$$

- $y=\text{True class:1}$ and $p: 0.3$
 $-(1 * \log(0.3) + (1 - 1) * \log(0.7)) = 0.52$
- $y=\text{True class:1}$ and $p: 0.8$
 $-(1 * \log(0.75) + (1 - 1) * \log(0.25)) = 0.1$

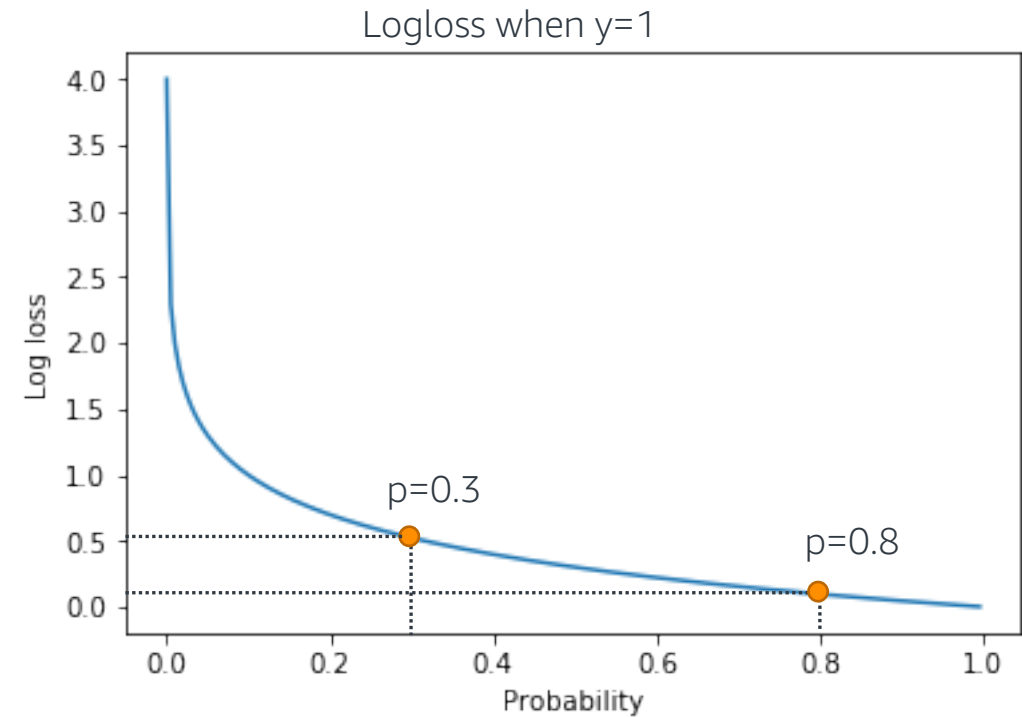


Log-loss (Binary Cross-entropy)

Example: Let's calculate the log-loss for these situations below

$$\text{Logloss} = -(y * \log(p) + (1 - y) * \log(1 - p))$$

- $y=\text{True class:1}$ and $p: 0.3$
 $-(1 * \log(0.3) + (1 - 1) * \log(0.7)) = 0.52$
- $y=\text{True class:1}$ and $p: 0.8$
 $-(1 * \log(0.75) + (1 - 1) * \log(0.25)) = 0.1$



Log-loss (Binary Cross-entropy)

Example: Let's calculate the log-loss for these situations below

$$\text{Logloss} = -(y * \log(p) + (1 - y) * \log(1 - p))$$


- $y=\text{True class:1}$ and $p: 0.3$

$$-(1 * \log(0.3) + (1 - 1) * \log(0.7)) = \boxed{0.52}$$

- $y=\text{True class:1}$ and $p: 0.8$

$$-(1 * \log(0.75) + (1 - 1) * \log(0.25)) = \boxed{0.1}$$

Better prediction gives smaller loss



Gradient Descent - Fitting a model

- Let's have our logistic regression model:

$$y = \text{sigmoid}(w_0 + w_1x_1 + w_2x_2 + \dots + w_qx_q)$$

where w_0, w_1, \dots, w_m : coefficients and x_0, x_1, \dots, x_m : input variables

- Minimize the cost function: **Log-loss**:

$$\text{Logloss} = \sum_{i=1}^n -(y_i * \log(p_i) + (1 - y_i) * \log(1 - p_i))$$

where


i : index, n : # number of records, y_i : True label, p_i : Probability of prediction

- Iteratively update coefficients with Gradient Descent:

$$w_{\text{new}} = w_{\text{current}} - \text{step_size} * \text{gradient}$$

Topics for today

- Optimization
- Regression
- **Boosting**
- Neural Networks
- MxNet
- Final Project



Adaptive Boosting (AdaBoost)

Main idea of boosting: Instead of learning a single model, learn multiple weak models that are good at different parts of the data.

It works as below:

- Output of the system (classification) is a weighted sum of individual models.
- We will first use a model using our training data. Then, create a second model that tries to correct initially misclassified samples.
- Add more models until all data correctly classified or we reach the max. number of iterations.


Adaptive Boosting (AdaBoost)

On each iteration t :


- Learn a weak model h_t and calculate its error rate ϵ_t
- Assign the strength of this weak model: α_t
- Assign a weight $D(i)$ to each data point i that is correlated to how incorrectly classified (We will use ϵ_t and α_t).

Final classifier: $H(X) = \text{sign}(\sum_{t=1}^T \alpha_t h_t(X))$


AdaBoost – Example




AdaBoost – Example




AdaBoost – Example




AdaBoost – Example




AdaBoost – Example



AdaBoost – Example









AdaBoost – Example




AdaBoost – Example

We also calculate each model's strength: α based on its correct predictions



AdaBoost – Example

Final classifier: $H(X) = \text{sign}(\sum_{t=1}^3 \alpha_t h_t(X))$



Topics for today

- Optimization
- Regression
- Boosting
- Neural Networks
- MxNet
- Final Project




Why Neural Networks?

- In the last 5-10 years, neural networks achieved **state-of-the art results** in many different machine learning areas.
- Neural networks can automatically **extract useful features** from input data. Traditionally feature engineering has been the most laborious and unprincipled part of building ML models
- When using neural networks, we should be careful about **model complexity** and make sure we have enough amount of data for training.


Real Neural Networks

- **Neuron** = cell that processes and transmits information through electrical and chemical signals
- **Synapses** = connections between neurons transmitting signals
- **Dendrites** = thin structures forming a complex tree receiving signals
- **Axon** = single output sending signals
- Neurons are **electrically excitable**, maintaining voltage gradients
- If the voltage changes by a large enough amount, an **electrochemical pulse** is generated and travels rapidly along the axon and activates synaptic connections with other cells when it arrives.
- Neurons form complex network structures, the human brains has 10^{11} neurons.




Artificial Neural Networks

- Has many processing units (neurons)
- Received numbers from many other neurons (synapses)
- Combines numbers into weighted sum (dendrites)
- Has single outgoing number (axon)
- Applies function to decide how strong of a signal to send (pulse)




Artificial Neural Networks

Perceptron: A single neuron within a connected set of neurons.

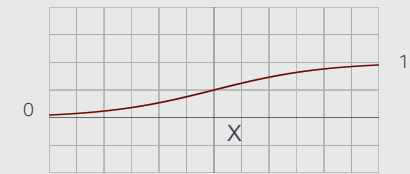
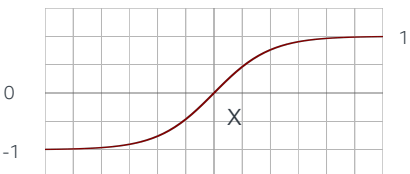



Multi Layer Perceptron




- A neural network consisting of input, hidden and output layers.
- Each layer is connected to the next layer.
- Except for the input units, each neuron has an activation function.

Activation Functions

Name	Plot	Function	Description
Logistic (sigmoid)		$f(x) = \frac{1}{1+e^{-x}}$	The most common activation function. Squashes input to $[0,1]$
Hyperbolic tangent (tanh)		$f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$	Squashes input to $[-1, 1]$
Rectified Linear Unit (ReLU)		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	Popular act. Function. Anything less than 0, results in zero activation

Derivatives of these functions are also important when we use gradient descent.


Forward Pass



Assume we have a network with:

- 2 inputs: $i_1=0.05$ and $i_2=0.1$
- 1 output (Binary classification: 0 or 1)
- 1 Hidden layer
- All neurons have sigmoid activation function

Forward Pass



$$net_{h1} = w_1 * i_1 + w_2 * i_2$$

$$net_{h1} = 0.15 * 0.05 + 0.2 * 0.1 = 0.0275$$


$$out_{h1} = \frac{1}{1 + e^{-net_{h1}}} = \frac{1}{1 + e^{-0.0275}} = 0.507$$

$$out_{h2} = 0.511$$

Sigmoid function

$$f(x) = \frac{1}{1 + e^{-x}}$$

Forward Pass



$$net_{o1} = w_5 * out_{h1} + w_6 * out_{h2}$$

$$net_{o1} = 0.4 * 0.507 + 0.45 * 0.511$$


$$out_{o1} = \frac{1}{1 + e^{-net_{o1}}} = \frac{1}{1 + e^{-0.432}} = 0.61$$

For binary classification, we would classify this as class 1 ($0.61 > 0.5$)

Output Function

- Binary classification: Sigmoid
 - Outputs $P(\text{target class}|\mathbf{x})$ in $[0,1]$
 - Last layer is Logistic Regression of output of previous layers
- Multi-class classification: Softmax
 - Still want P for each class output in $[0,1]$
 - Want sum of output to be 1 (probability distribution)
 - Training drives value for target class up, others down.
- Regression: Output activation can be linear or relu

$$\sigma(z) = \frac{1}{1+\exp(-z)}$$



$$\text{softmax}(\mathbf{z})_i = \frac{\exp(z_i)}{\sum_j \exp(z_j)}.$$

Cost Functions

Binary classification: Cross entropy for logistic

$$C = -\frac{1}{n} \sum_{\text{examples}} y \ln p + (1 - y) \ln(1 - p)$$

Multi-class classification: Cross entropy for Softmax

$$C = -\frac{1}{n} \sum_{\text{examples}} \sum_{\text{classes}} y_j \ln p_j$$

Regression: Mean Squared Error: $C = \frac{1}{n} \sum_{\text{examples}} (y - p)^2$

Notation for Classification

- n training examples
- j classes
- p = prediction (probability)
- y = true class (1 for yes, 0 for no)

Notation for Regression

- n training examples
- p = prediction (numeric)
- y = true value


Training Neural Networks

- During the training of neural networks, we can update the weights in the network by applying the gradient descent method.
- Cost function is selected accr. to problem: Binary, multi-class classification or regression.
- Weight update formula:

$$w_n: w_n - \text{learning_rate} * \left(\frac{\partial E}{\partial w_n} \right)$$


E: Error

Gradient of w_n

An orange arrow points from the term $\left(\frac{\partial E}{\partial w_n} \right)$ in the formula to the text "Gradient of w_n ".

Dropout

- Regularization technique to prevent overfitting.
- Randomly removes some nodes and their connections during the training.



Topics for today

- Optimization
- Regression
- Boosting
- Neural Networks
- **MXNet**
- Final Project



MXNet



- Open source Deep Learning Library to train and deploy neural networks.
- With the **Gluon** interface, we can define and train neural networks easily.
- We will go over the tutorials below:
 - **Mxnet-Ndarrays-Autograd intro:**
<https://eider.corp.amazon.com/sazaracs/notebook/NBM9GVVWI98P>
 - **Building Neural Network with MXNet:**
<https://eider.corp.amazon.com/sazaracs/notebook/NBLQY2I99MEG>




Hands-on exercise

- In this exercise we will work with an internal Amazon Dataset.
- We will do the following tasks:
 - Exploratory Data Analysis
 - Splitting dataset into training and test sets.
 - Fit a Neural Network
 - Check the performance metrics on test set.
- Open the notebook below:

<https://eider.corp.amazon.com/sazaracs/notebook/NB8A43NS520E>



Dive into Deep Learning



Dive into Deep Learning

An **interactive** deep learning book with code, math, and discussions, based on the **NumPy** interface.

E-book on deep learning by Amazon Scientists, available here: <https://d2l.ai>

Related chapters:

Chapters 3: Linear Neural Networks: https://d2l.ai/chapter_linear-networks/index.html

Chapters 4: Multilayer Perceptrons: https://d2l.ai/chapter_multilayer-perceptrons/index.html

AutoML

AutoML helps automating some of the tasks related to ML model development and training such as:

- Preprocessing and cleaning data
- Feature selection
- ML model selection
- Hyper-parameter optimization

Auto LUON AutoML

- Open source AutoML Toolkit (AMLT) created by Amazon AI
- With **AutoGluon**, state of the art ML results can be achieved in a few lines of Python code.

```
>>> from autogluon import TabularPrediction as task
>>> predictor = task.fit(train_data=task.Dataset(file_path=TRAIN_DATA.csv), label=COLUMN_NAME)
>>> predictions = predictor.predict(task.Dataset(file_path=TEST_DATA.csv))
```


Auto G LUON AutoML

- Easy to Use – Built-in Applications


Tabular Prediction

system 1 status on hold						
	Region	Product	Subproduct	Unit	Plan	Actual
101	United States	Alcohol	New York	10	1000	1000
102	United States	Alcohol	New York	10	1000	1000
103	United States	Alcohol	New York	10	1000	1000
104	United States	Alcohol	New York	10	1000	1000
105	United States	Alcohol	New York	10	1000	1000
106	United States	Alcohol	New York	10	1000	1000
107	United States	Alcohol	New York	10	1000	1000
108	United States	Alcohol	New York	10	1000	1000
109	United States	Alcohol	New York	10	1000	1000
110	United States	Alcohol	New York	10	1000	1000
111	United States	Alcohol	New York	10	1000	1000
112	United States	Alcohol	New York	10	1000	1000
113	United States	Alcohol	New York	10	1000	1000
114	United States	Alcohol	New York	10	1000	1000
115	United States	Alcohol	New York	10	1000	1000
116	United States	Alcohol	New York	10	1000	1000
117	United States	Alcohol	New York	10	1000	1000
118	United States	Alcohol	New York	10	1000	1000
119	United States	Alcohol	New York	10	1000	1000
120	United States	Alcohol	New York	10	1000	1000

Image Classification



Object Detection



Text Classification



Auto LUON AutoML

- AutoGluon is able to produce top solutions to many real-world business problems. AutoGluon was utilized to create:
 - The first place solution to the MLA NLP I Hazmat Challenge
 - First place against over 1,000 competitors and 10,000 manual submissions over the course of 12 months.
 - The first place solution to the MLE Ops Tech IT Challenge
 - The first place solution to the SCOT Grand Challenge 2018
 - Solution was able to determine if two products are identical better than a human evaluator.
 - Used as the production solution to the product matching problem within Amazon.
 - The first place solution to the SCOT Grand Challenge 2019
 - Solution approximately halved the predictive error of Amazon.com shipment costs worldwide.
 - The existing production solution to Amazon Grocery (F3) sales forecasting
- pip install to get started: <https://github.com/awsmlabs/autogluon>

Topics for today

- Optimization
- Regression
- Boosting
- Neural Networks
- MXNet
- Final Project



Final Project

Product substitute: Given products (A, B), predict whether B is a substitute for A

- We say that B is a "substitute" for A if a customer would buy B in place of A -- say, if A were out of stock.
- **The goal** of this project is to predict a substitute relationship between pairs of products.
- **Link:** <https://leaderboard.corp.amazon.com/tasks/478>

Final Project Walkthrough – Day 3

- See the provided project walkthrough below to get started with this project.
- You can use Boosted Trees and Neural Networks this time.
- Complete the following notebook and submit your results
<https://eider.corp.amazon.com/sazaracs/notebook/NBLR85NVZ4V4>



Notebooks for day 3

- MXNet – NDarrays intro:
<https://eider.corp.amazon.com/sazaracs/notebook/NBM9GVVWI98P>
- Neural Networks intro:
<https://eider.corp.amazon.com/sazaracs/notebook/NBLQY2I99MEG>
- Neural Networks full example:
<https://eider.corp.amazon.com/sazaracs/notebook/NB8A43NS520E>
- Final project walkthrough for day 3:
<https://eider.corp.amazon.com/sazaracs/notebook/NBLR85NVZ4V4>